

# challenge-1

March 15, 2024

## 1 C1.1 Evaluation on Object Detection

### 1.1 Team: *Paradox*

### 1.2 Teammates: ‘kjamunap’ , ‘sriramma’ , ‘leonli’

### 1.3 Introduction

- Task: Evaluate the average precision of the data collected from SafeBench.
- Submission: the plotted P-R curve, as well as the average precision at different IoU threshold level.

```
[51]: import joblib
import numpy as np
import pandas as pd
from matplotlib import pyplot as plt
```

```
[52]: !git clone https://github.com/HenryLHH/24784-c1-data.git
!cp /content/24784-c1-data/results.pkl .
!rm -r /content/24784-c1-data/
```

```
Cloning into '24784-c1-data'...
remote: Enumerating objects: 4, done.
remote: Counting objects: 100% (4/4), done.
remote: Compressing objects: 100% (4/4), done.
remote: Total 4 (delta 0), reused 4 (delta 0), pack-reused 0
Receiving objects: 100% (4/4), 31.64 KiB | 31.64 MiB/s, done.
```

#### 1.3.1 Load the Detection Results from SafeBench\_v2

```
[53]: inputs = joblib.load('results.pkl') # move the files to your current folder
inputs.keys() # dict_keys(['image_id', 'predicted_class', 'ground_truth_bbox', 'predicted_bbox', 'conf_scores', 'num_labels'])

[53]: dict_keys(['image_id', 'predicted_class', 'ground_truth_bbox', 'predicted_bbox', 'conf_scores', 'num_labels'])
```

### 1.3.2 Compute the IoU

```
[54]: def box_area(box):
    """ Compute the box area, given all the vertices
    # Arguments
        box: (4, N) ndarray
    # Returns
        areas: (N, ) ndarray
    """
    assert box.shape[0] == 4

    width = np.abs(box[0] - box[2])
    height = np.abs(box[1] - box[3])

    areas = width * height

    # TODO, compute the rectangle areas based on the vertices input
    return areas

def box_iou(box1, box2):
    """
        Compute the iou between box1 and box2
        ONLY consider the SINGLE ground truth, which is a simplified case
        box1: (N, 4) ndarray
        box2: (N, 4) ndarray
        return: (N, ) iou scores
    """
    eps = 1e-7

    assert box1.ndim == 2 and box2.ndim == 2, "Inputs must be 2-dimensional arrays"
    assert box1.shape[1] == 4 and box2.shape[1] == 4, "Each box should have 4 coordinates"

    a1, a2 = np.split(box1, 2, axis=1)
    b1, b2 = np.split(box2, 2, axis=1)

    # Calculate the coordinates of intersection rectangle
    x1 = np.maximum(a1[:, 0], b1[:, 0])
    y1 = np.maximum(a1[:, 1], b1[:, 1])
    x2 = np.minimum(a2[:, 0], b2[:, 0])
    y2 = np.minimum(a2[:, 1], b2[:, 1])

    # Compute the width and height of intersection rectangle
    inter_width = np.maximum(x2 - x1, 0)
    inter_height = np.maximum(y2 - y1, 0)
```

```

# Compute the area of intersection
inter = inter_width * inter_height

# Compute the area of union
box1_area = (box1[:, 2] - box1[:, 0]) * (box1[:, 3] - box1[:, 1])
box2_area = (box2[:, 2] - box2[:, 0]) * (box2[:, 3] - box2[:, 1])
union = box1_area + box2_area - inter

return inter / (union + eps)

```

1. In the `box_area` function, the area of each bounding box is calculated based on its vertices. We have a numpy array `box` of shape  $(4, N)$ , where  $N$  is the number of bounding boxes.
2. Then we calculate the width and height of each box by taking the differences between the x and y coordinates of opposite vertices, and then multiply these values to get the area.
3. The `box_iou` function, computes IoU scores between pairs of bounding boxes.
4. We take two numpy arrays `box1` and `box2`, each of shape  $(N, 4)$ , representing  $N$  bounding boxes with four coordinates  $(x_1, y_1, x_2, y_2)$ .
5. Calculating the coordinates of the intersection rectangle by finding maximum x and y coordinates of top-left corner and minimum x and y coordinates of bottom-right corner.
6. Then we find width and height of the intersection rectangle and calculate the area of intersection.
7. The area of the union of the two bounding boxes and divides the area of intersection by the area of the union to obtain the IoU score.
8. The IoU loss is defined as the complement of the IoU scores.

```
[55]: # TODO: get the box iou scores from the function you implemented above
iou_scores = box_iou(inputs['ground_truth_bbox'], inputs['predicted_bbox'])
# Compute IoU loss
iou_loss = 1.0 - iou_scores

print("IoU Loss values:")
print(iou_loss)
```

IoU Loss values:  
[1. 1. 1. ... 1. 1. 1.]

```
[56]: average_loss = np.mean(iou_loss)

# Print the average loss
print("Average IoU Loss:", average_loss)
```

Average IoU Loss: 0.9238258051395144

```
[57]: # Build your dataframe from the dictionary
input_dict = {
    'image_id': inputs['image_id'],
    'predicted_class': inputs['predicted_class'],
    'conf_scores': inputs['conf_scores'],
```

```

        'iou_scores': iou_scores,
    }
df = pd.DataFrame.from_dict(input_dict)
df

```

[57]:

	image_id	predicted_class	conf_scores	iou_scores
0	0	None	-1.000000	0.000000
1	1	car	0.671808	0.000000
2	1	car	0.563249	0.000000
3	1	car	0.452749	0.80052
4	1	car	0.428004	0.000000
...	...	...	...	...
1631	91	train	0.219339	0.000000
1632	91	bus	0.173378	0.000000
1633	91	person	0.101061	0.000000
1634	91	sink	0.088729	0.000000
1635	91	dining table	0.073306	0.000000

[1636 rows x 4 columns]

[58]:

```

df = df[df.conf_scores >= 0] # drop all the frames without detection (conf_scores = -1)

# get all the detection results of stop signs
df_stopsign = df.loc[(df.predicted_class=='stop sign')]
df_stopsign

```

[58]:

	image_id	predicted_class	conf_scores	iou_scores
160	7	stop sign	0.071092	0.812147
345	15	stop sign	0.052572	0.823549
712	33	stop sign	0.110136	0.814429
731	34	stop sign	0.159796	0.844884
772	36	stop sign	0.128940	0.850620
...	...	...	...	...
1606	88	stop sign	0.119673	0.025573
1610	89	stop sign	0.997623	0.932374
1622	90	stop sign	0.995437	0.923847
1628	90	stop sign	0.052746	0.032057
1629	91	stop sign	0.995099	0.614191

[66 rows x 4 columns]

## 1.4 Compute the Average Precision

```
[59]: def interp_ap(recall, precision, method = 'interp'):
    """ Compute the average precision, given the recall and precision curves
    # Arguments
        recall: The recall curve (list)
        precision: The precision curve (list),
        methods: 'continuous', 'interp'
    # Returns
        Average precision, precision curve, recall curve
    """

    # TODO: Append sentinel values to beginning and end
    # Recall should start with 0.0 and end with 1.0
    # Precision should start with 1.0 and end with 0.0

    appended_recall = np.concatenate(([0.0], recall, [1.0]))
    appended_prec_input = np.concatenate(([1.0], precision, [0.0]))

    # Compute the precision envelope
    appended_prec = np.flip(np.maximum.accumulate(np.
        flip(appended_prec_input)))  # get the  $p(r) = \max_{r' \leq r} p(r')$ 

    # Integrate area under curve
    if method == 'interp':
        x = np.linspace(0, 1, 101)  # 101-point interp (COCO)
        ap = np.trapz(np.interp(x, appended_recall, appended_prec), x)  # get the average precision based on the areas under interpolated curves
    else:  # 'continuous'
        i = np.where(appended_recall[1:] != appended_recall[:-1])[0]  # points where x axis (recall) changes
        ap = np.sum((appended_recall[i + 1] - appended_recall[i]) * appended_prec[i + 1])  # area under curve

    return ap, appended_prec_input, appended_prec, appended_recall
```

```
[60]: def compute_ap(df, num_gt, iou_thres):
    """ Compute the average precision, given the recall and precision curves
    # Arguments
        df: DataFrame inputs containing predicted classes, iou_scores, and confidence
        num_gt: The precision curve (list),
        iou_thres: IoU threshold of True Positive (TP) detection
    # Returns
        Average precision, precision curve, recall curve
    """

    # Create a copy of the DataFrame
    df = df.copy()

    # Sort by confidence in descending order
    df.sort_values('confidence', ascending=False, inplace=True)

    # Initialize lists to store precision and recall
    precision = []
    recall = []

    # Initialize variables for cumulative sum
    total_tp = 0
    total_fp = 0
    total_gt = num_gt

    # Loop through each row in the sorted DataFrame
    for _, row in df.iterrows():
        # If the row is a True Positive (TP)
        if row['iou_scores'] >= iou_thres:
            total_tp += 1
            # Add the current row to the precision curve
            precision.append(total_tp / total_gt)
            # Add the current row to the recall curve
            recall.append(total_tp / total_gt)
        else:
            # Add the current row to the precision curve
            precision.append(total_tp / total_gt)
            # Add the current row to the recall curve
            recall.append(total_tp / total_gt)

    # Compute the average precision using the interp_ap function
    ap, _, _, _ = interp_ap(recall, precision, method='interp')

    return ap, precision, recall
```

```

df = df.sort_values(by='conf_scores', ascending=False)      # TODO sort all
    ↵the data with confidence scores 'conf_scores'
tp_fp = np.arange(1, len(df)+1, 1)
tp = ((df['iou_scores'] >= iou_thres) & (df['predicted_class'] ==
    ↵'stopsign')).cumsum().values    # get the true positive sets
precision = tp / tp_fp        # array (N, )
recall = tp / num_gt          # array (N, )

ap, prec_raw, prec, recall = interp_ap(recall, precision)      # get
    ↵the AP and returned curve for plot

return ap, prec_raw, prec, recall

```

## 1.5 Visualization

```
[61]: def plot_pr_curves(recall, prec_raw, prec):
    """ Visualize the P-R curves
        Visualize the areas under P-R curves
    """
    plt.figure()
    plt.plot(recall, prec_raw)
    plt.plot(recall, prec)
    plt.fill_between(recall, np.zeros_like(recall), prec, color='orange',
    ↵alpha=0.2)
    plt.legend(['Raw curve', 'Precision Envelope', 'Average Precision'])
    plt.title('P-R Curve')
    plt.xlabel('Recall')
    plt.ylabel('Precision')
```

## 1.6 Execute the AP calculation

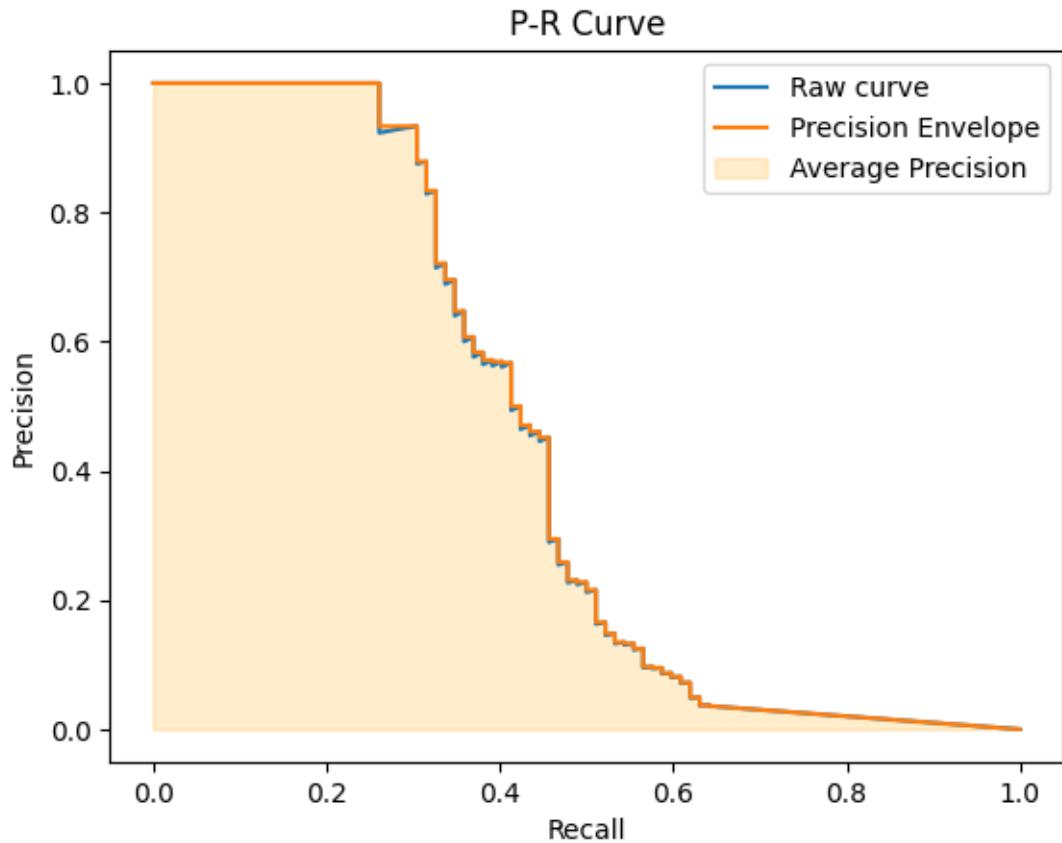
```
[62]: # Traverse over IoU threshold, get AP@[0.5:0.1:0.9]

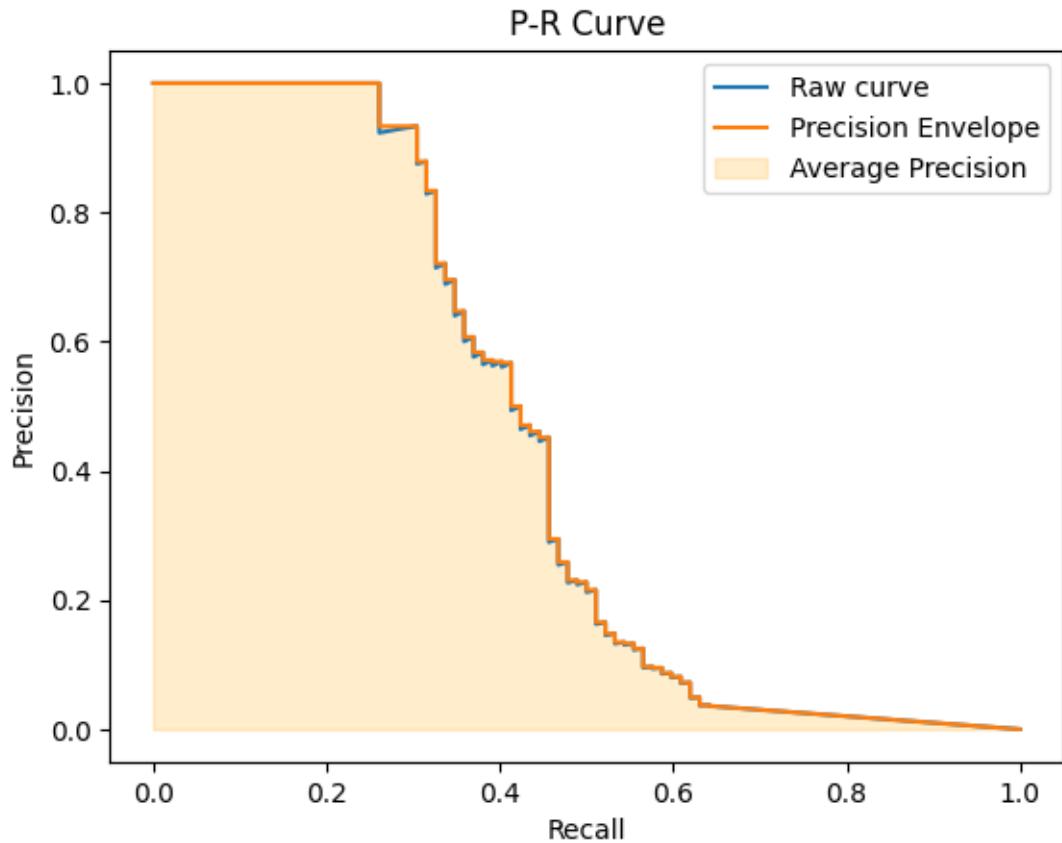
for iou_thres in np.arange(0.5, 1.0, 0.1):
    ap, prec_raw, prec, recall = compute_ap(df, inputs['num_labels'], iou_thres)
    print('IoU Threshold: {:.2f} | AP@{:.2f}'.format(iou_thres),
    ↵format(ap))
    plot_pr_curves(recall, prec_raw, prec)
```

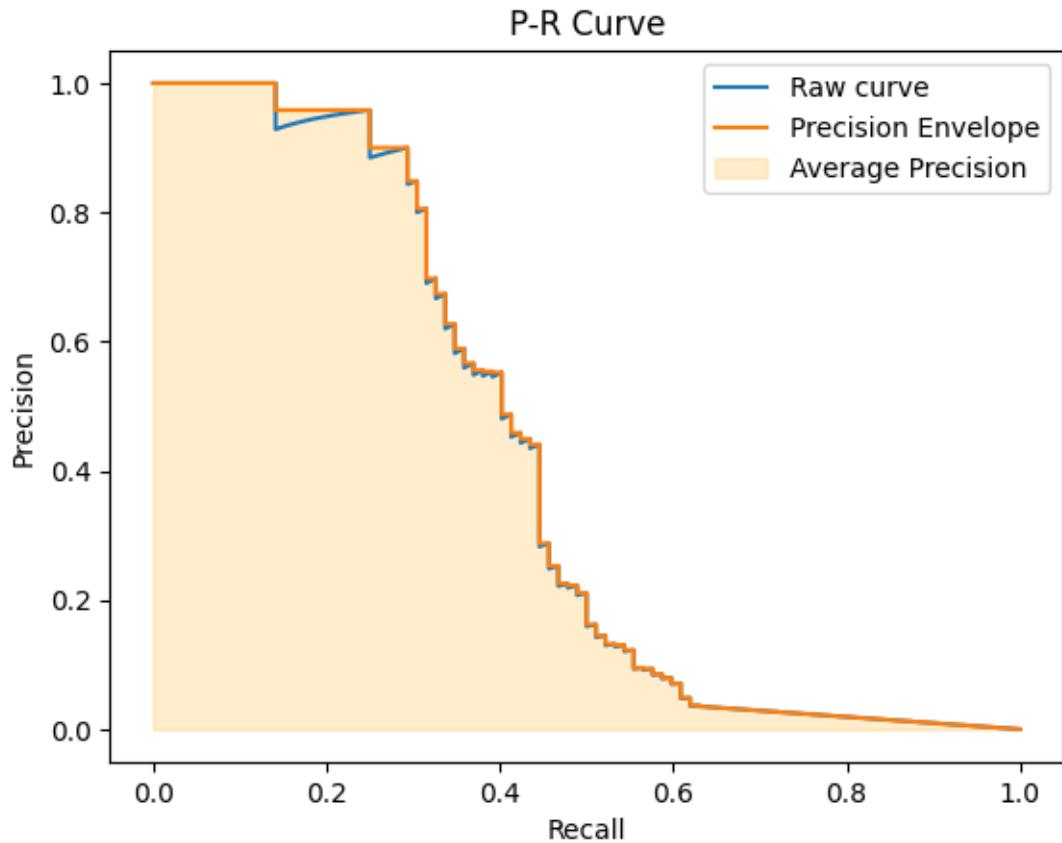
```

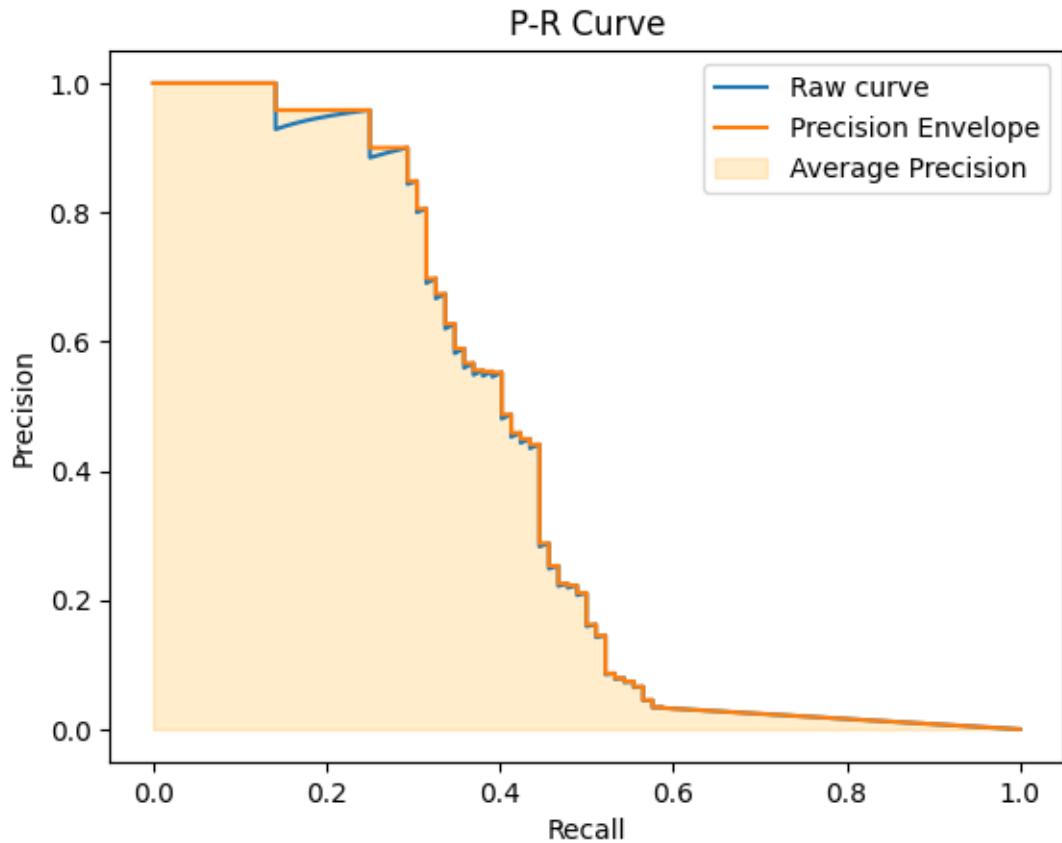
IoU Threshold: 0.50 | AP@0.50 0.4272559074361477
IoU Threshold: 0.60 | AP@0.60 0.4272559074361477
IoU Threshold: 0.70 | AP@0.70 0.40686292416368364
IoU Threshold: 0.80 | AP@0.80 0.4018855902817313
IoU Threshold: 0.90 | AP@0.90 0.3466670309329718

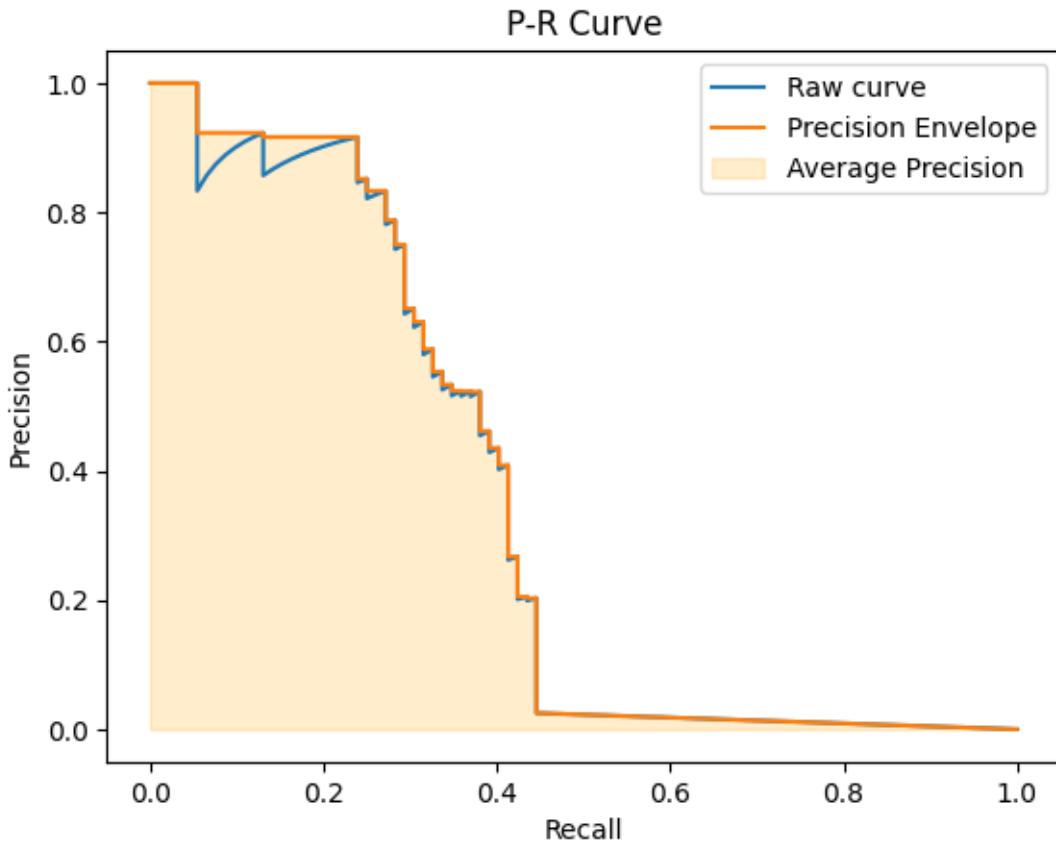
```











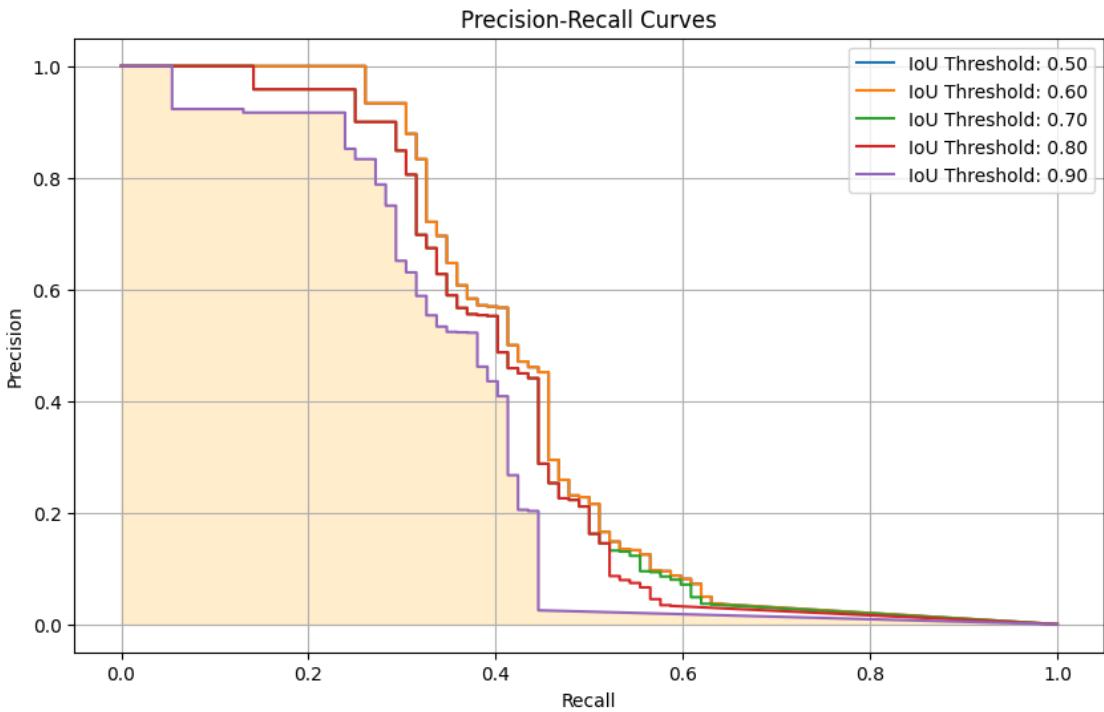
### *Precision-Recall curves with envelope curves*

```
[63]: for iou_thres in iou_thresholds:
    ap, _, _, _ = compute_ap(df, inputs['num_labels'], iou_thres)
    ap_values.append(ap)

plt.figure(figsize=(10, 6))

for iou_thres in iou_thresholds:
    _, prec_raw, prec, recall = compute_ap(df, inputs['num_labels'], iou_thres)
    plt.plot(recall, prec, label=f'IoU Threshold: {iou_thres:.2f}')

plt.fill_between(recall, np.zeros_like(recall), prec, color='orange', alpha=0.2)
plt.legend()
plt.title('Precision-Recall Curves')
plt.xlabel('Recall')
plt.ylabel('Precision')
plt.grid(True)
plt.show()
```



```
[64]: iou_thresholds = np.arange(0.5, 1.0, 0.1)
ap_values = []

for iou_thres in iou_thresholds:
    ap, _, _, _ = compute_ap(df, inputs['num_labels'], iou_thres)
    ap_values.append(ap)

results_df = pd.DataFrame({
    '|IoU-Threshold | AP@IoU|': [f'| {iou:.2f} | {ap:.4f} |'
        for iou, ap in zip(iou_thresholds, ap_values)]
})
print(results_df.to_string(index=False))
```

IoU-Threshold	AP@IoU
0.50	0.4273
0.60	0.4273
0.70	0.4069
0.80	0.4019
0.90	0.3467

IoU Threshold	AP@IoU
0.50	0.4273
0.60	0.4273

<b>IoU Threshold</b>	<b>AP@IoU</b>
0.70	0.4069
0.80	0.4019
0.90	0.3467

- The trend observed as the IoU threshold increases, a gradual decrease in the Average Precision (AP). From the results, as the IoU threshold increases from 0.50 to 0.90, the AP decreases from 0.4273 to 0.3467.
- As the IoU threshold increases, the number of true positives decreases.
- This is because only predicted bounding boxes with a higher overlap with the ground truth bounding boxes are counted as true positives.
- Bounding boxes with lower overlaps are considered false positives.
- With fewer true positives and large number of false positives, both precision and recall decrease as the IoU threshold increases.
- We know that precision usually decreases at a faster rate than recall, this affect we can see in a decrease in AP.
- Higher IoU thresholds require more precise localization of objects by the model.
- This can be challenging, especially in cases where objects have complex shapes.